

Scheduler Comparison with 95% CI — ShareGPT
LLaMA-3-8B, 1 A100 80GB GPU (n=3 runs)

